

Homework 6

Farnoosh Moshirfatemi

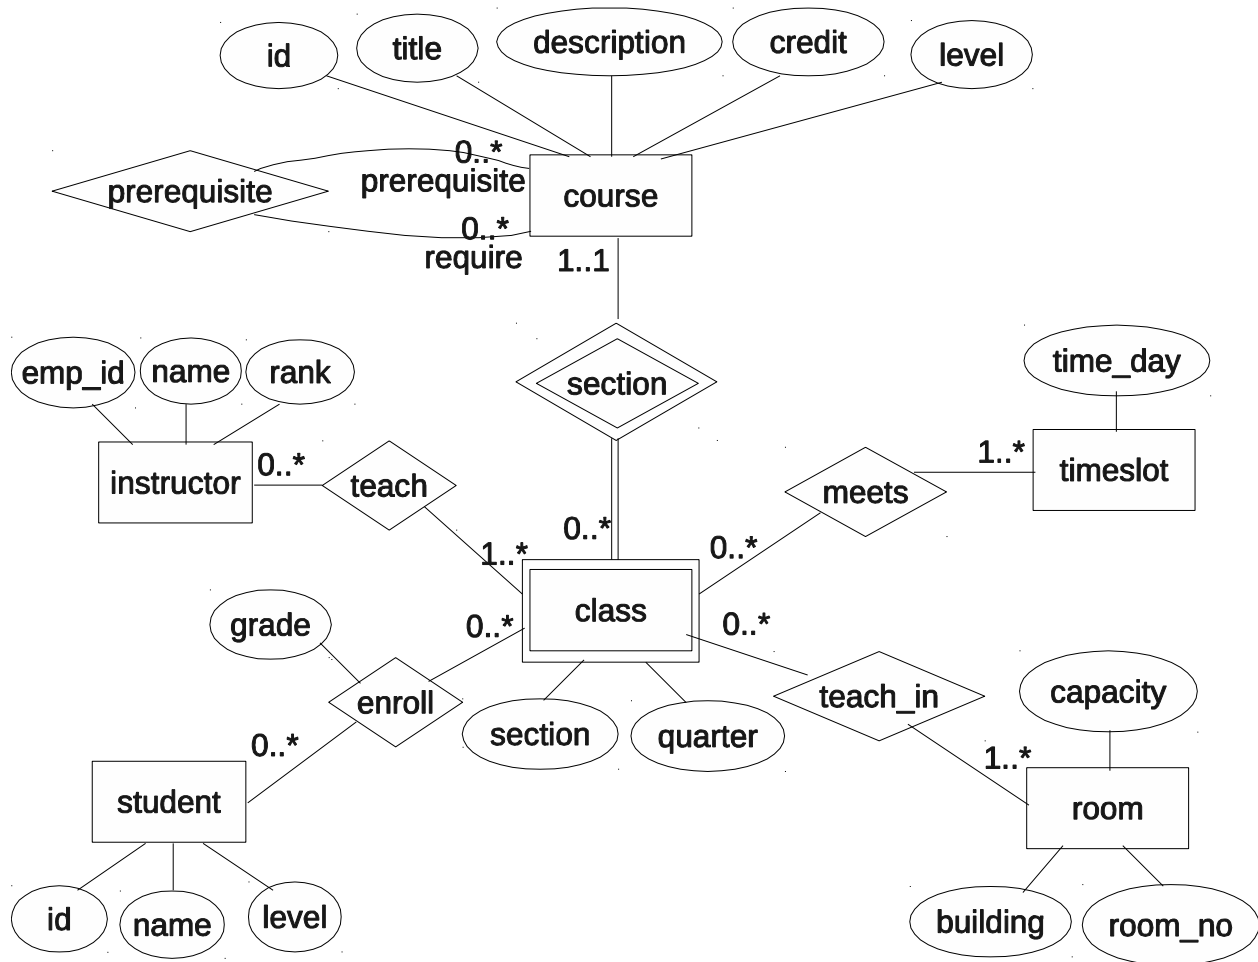
Dejun Qian

Question 1:

Database design problem

Answer:

The ER diagram of the database is shown below.



By translating the ER diagram, we get the following set of tables. The tables are described using create table statements.

Table 1. Create table statements for the database

table name	create table statement
course	<pre>CREATE TABLE course (id CHARACTER(22) PRIMARY KEY, title CHARACTER (20) NOT NULL, description CHARACTER (20), credit INTEGER, level CHARACTER(2));</pre>
prerequisit	<pre>CREATE TABLE prerequisite (course_id CHARACTER(22) REFERENCES course(id), precourse_id CHARACTER(22) REFERENCES course(id), PRIMARY KEY(course_id, precourse_id));</pre>
class	<pre>CREATE TABLE class (id CHARACTER(22) REFERENCES course(id), section INTEGER, quarter CHARACTER (22), PRIMARY KEY(id, section, quarter));</pre>
instructor	<pre>CREATE TABLE instructor (emp_id INTEGER PRIMARY KEY, name CHARACTER(20), rank INTEGER);</pre>
teach	<pre>CREATE TABLE teach (id CHARACTER(22), section INTEGER , quarter CHARACTER (22), emp_id INTEGER REFERENCES instructor(emp_id), PRIMARY KEY(id, section, quarter, emp_id), FOREIGN KEY(id, section, quarter) REFERENCES class(id, section, quarter));</pre>
student	<pre>CREATE TABLE student (id INTEGER PRIMARY KEY, name CHARACTER(22), level CHARACTER(1));</pre>

enroll	<pre> CREATE TABLE enroll (course_id CHARACTER(22), section INTEGER, quarter CHARACTER (22), student_id INTEGER REFERENCES student(id), grade CHARACTER(1), PRIMARY KEY(course_id, section, quarter, student_id), FOREIGN KEY(course_id, section, quarter) REFERENCES class(id, section, quarter)); </pre>
room	<pre> CREATE TABLE room (building CHARACTER(20), room_no INTEGER, capacity INTEGER, PRIMARY KEY(building, room_no)); </pre>
teach_in	<pre> CREATE TABLE teach_in (id CHARACTER(22), section INTEGER, quarter CHARACTER (22), room_no INTEGER, building CHARACTER(20), PRIMARY KEY(id, section, quarter, room_no, building), FOREIGN KEY(room_no, building) REFERENCES room(room_no, building), FOREIGN KEY(id, section, quarter) REFERENCES class(id, section, quarter)); </pre>
timeslot	<pre> CREATE TABLE timeslot (time_day CHARACTER(20) PRIMARY KEY); </pre>
meets	<pre> CREATE TABLE meets (id CHARACTER(22), section INTEGER, quarter CHARACTER (22), time_day CHARACTER(22) REFERENCES timeslot(time_day), PRIMARY KEY(id, section, quarter, time_day), FOREIGN KEY(id, section, quarter) REFERENCES class(id, section, quarter)); </pre>

For the constraints listed, none of them could be fully supported by a foreign key constraint. Since we wanted the constraints to be enforced at all time, and not only before or after some triggering event, we implemented all of the constraints with assertions and not triggers. We preferred to use SQL statement instead of embedded SQL program, because this way we don't to write any C or Java code. SQL statement is at higher level as C or Java.

Therefore, we need to create assertions for each of them.

The results are shown below.

1-A room cannot be booked for more than one class (section) at the same time slot.

For this constraint, we wrote an assertion to check there are no rows in the table where the number of reservation for a room and a time slot is bigger than 1. The constraint is shown below.

```
CREATE ASSERTION question1 CHECK(NOT EXISTS
(SELECT *
FROM ( SELECT count(*) AS C
FROM timeslot, meets, teach_in, room
WHERE timeslot.time_day = meets.time_day AND
meets.id = teach_in.id AND
meets.section = teach_in.section AND
meets.quarter = teach_in.quarter AND
teach_in.building = room.building AND
teach_in.room_no = room.room_no
group by timeslot.time_day, room.building, room.room_no) x
WHERE C > 1
)
```

2-The number of students enrolled in a class cannot exceed the capacity of the room assigned to the class.

Similar to the previous constraint, we wrote an assertion to check there are no rows in the table where the number of students enrolled in a class exceeds the capacity of the room assigned to the class. The constraint is shown below.

```
CREATE ASSERTION question1 CHECK(NOT EXISTS
(SELECT enroll.course_id, enroll.quarter, enroll.section, room.building, room.room_no
FROM student, enroll, teach_in, room
WHERE student.id = enroll.student_id AND
enroll.course_id = teach_in.id AND
enroll.section = teach_in.section AND
enroll.quarter = teach_in.quarter AND
teach_in.building = room.building AND
teach_in.room_no = room.room_no
```

```
group by enroll.course_id, enroll.quarter, enroll.section, room.building, room.room_no
having count(*)>room.capacity)
)
```

3- Students may not take more than 21 credits in any one quarter. (PSU rules may be more flexible.)

For this constraint, we wrote an assertion to check there are no rows in the table where students take more than 21 credits in any one quarter. The SQL statement is shown below.

```
CREATE ASSERTION question1 CHECK(NOT EXISTS
(SELECT class.quarter, enroll.student_id
FROM enroll, class, course
WHERE
enroll.course_id = class.id AND
enroll.section = class.section AND
enroll.quarter = class.quarter AND
class.id = course.id
group by class.quarter, enroll.student_id
having sum(course.credit)>21)
)
```

4-A student can't enroll in two sections at the same time slot in a quarter.

To enforce this constraint, we wrote an assertion to enforce there are no rows in the table where a student enrolls in more than one section at the same time slot in a quarter. The constraint is shown below.

```
CREATE ASSERTION question1 CHECK(NOT EXISTS
(SELECT count(*) AS C
FROM timeslot, meets, enroll, student
WHERE timeslot.time_day = meets.time_day AND
meets.id = enroll.course_id AND
meets.section = enroll.section AND
meets.quarter = enroll.quarter AND
enroll.student_id = student.id
group by student.id, timeslot.time_day, enroll.quarter
having count(*)>1)
)
```

5-A student can't enroll in two sections of the same class in a quarter.

This constraint is done by writing an assertion to enforce there are no rows in the table where a student enrolls in more than one section of the same class in a quarter. The SQL statement is shown below.

```
CREATE ASSERTION question1 CHECK(NOT EXISTS
(SELECT enroll.student_id, class.id, class.quarter
FROM enroll, class
WHERE class.id = enroll.course_id AND
```

```
class.section = enroll.section AND  
class.quarter = enroll.quarter  
group by enroll.student_id, class.id, class.quarter  
having count(*)>1  
)
```

6- A student can't enroll in a class unless they have earned a grade of C or higher in the prerequisite courses.

We enforce this constraint by writing an assertion to make sure there are no rows in the table where a student enrolls in a class when they have earned a grade less than C in the prerequisite courses. The following is the SQL statement for this constraint.

```
CREATE ASSERTION question1 CHECK(NOT EXISTS  
(SELECT *  
FROM enroll e1, class, prerequisite, enroll e2  
WHERE class.id = e1.course_id AND  
class.section = e1.section AND  
class.quarter = e1.quarter AND  
class.id = prerequisite.course_id AND  
prerequisite.precourse_id = e2.course_id AND  
e2.grade < 'C')  
)
```

7- A student can't take the same course more than three times.

This constraint is done by writing an assertion to enforce there are no rows in the table where a student takes the same course more than three times. The SQL statement is shown below.

```
CREATE ASSERTION question1 CHECK(NOT EXISTS  
(SELECT count(*)  
FROM enroll  
GROUP BY enroll.student_id, enroll.course_id  
having count(*)>3)  
)
```

Note: As it is shown, all of the constraints are written as assertion statements. Therefore, they are guaranteed to be true at all time

The following is the answer for the query questions.

Q1: Write a query that will show the class roster for each class.

The query statement is shown below.

```
SELECT enroll.course_id, enroll.section, enroll.quarter, student.id, student.name
FROM enroll, student
where enroll.student_id = student.id
GROUP BY enroll.course_id, enroll.section, enroll.quarter, student.name, student.id
ORDER BY enroll.course_id, enroll.section, enroll.quarter
```

Q2: Write a query that will show the transcript for each student.

The query statement is shown below.

```
SELECT student.id, enroll.course_id, enroll.section, enroll.quarter, enroll.grade
FROM enroll, student
where enroll.student_id = student.id
GROUP BY student.id, enroll.course_id, enroll.section, enroll.quarter, enroll.grade
```

Question 2:

Normalization problem

Answer:

1- Identify all of the non-trivial FDs that hold (based on the application).

mission_id -> mission_name

mission_id -> access_id

mission_id -> team_id

mission_id -> mission_status

mission_id -> team_name

mission_id -> meeting_frequency

team_id -> team_name

team_id -> meeting_frequency

2- Identify the key(s) for this table.

mission_id is the key for this new table.

3- Identify any “troublesome” FDs that prevent this table from being in BCNF.

A table is in BCNF if the left side of every nontrivial FD be a superkey. Therefore, in this table the troublesome FDs are:

team_id -> team_name

team_id -> meeting_frequency

4- Describe one insert anomaly, one update anomaly, and one delete anomaly that can arise with this table.

Insert anomaly:

If we want to insert different missions for one team, we need to know the descriptive information for that team (team_id, team_name, team_meeting_frequency) and make sure this information is in consistence with the other rows with the same team_id wich are already in the table.

Also, we have to insert the same information for all the rows that you want to be inserted with the same team_id which is providing the redundant information.

Also, if we want to insert a team information we can't until there is at least one mission assigned to that team.

Update anomaly:

If we want to change the name of a specific team, we have to change the name as many times as it is used in the rows of the table.

Delete anomaly:

If one team has only one mission, by deleting the mission the team information would also be deleted.

Also, if we want to delete a specific team from the table, we have to delete all the rows that have that team_id.

5- Given that the system designers have decided to use this table and knowing that there are redundancies, describe (in English) the triggers that would need to be implemented in order to correctly manage the redundancy. You want to make sure that a given piece of information (if it is represented redundantly) is always consistent. That is, you want to make sure that all of the copies of any redundant information have the same, most up-to-date value. You want to make sure that if information is deleted, all of the copies are deleted, etc.

Basically, we have to have triggers for all the mentioned anomalies.

Trigger on update:

If a team information (team_name or team_meeting_freq) wants to be updated in the table, the trigger should be executed after the update and it should execute the same update for all rows of the table with the same team_id.

Trigger on insert:

The trigger should be executed before the insertion and make sure that the desired information to be inserted are in consistence with the previous information in the table and if this is not the case avoid the insertion.

Trigger on delete:

In order to delete a team from the table, a trigger should be executed after the deletion and delete all other rows in the table which have the same team_id as the deleted team.